

```

import pandas as pd
import numpy as np
import os
import tensorflow as tf
import matplotlib.pyplot as plt
import seaborn as sns
import gdown
import zipfile
from sklearn.model_selection import train_test_split
from sklearn.metrics import classification_report
import cv2
from tensorflow.keras.utils import Sequence
import albumentations as A
from albumentations.core.transforms_interface import DualTransform
from tensorflow.keras.optimizers.schedules import ExponentialDecay
from tensorflow.keras.callbacks import ReduceLROnPlateau
from tensorflow.keras import layers, models, callbacks
from tensorflow.keras.regularizers import l2
from sklearn.utils.class_weight import compute_class_weight

```

```
print(tf.__version__)
```

```

2025-03-19 19:44:54.310968: E
external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:467] Unable to
register cuFFT factory: Attempting to register factory for plugin
cuFFT when one has already been registered
WARNING: All log messages before absl::InitializeLog() is called are
written to STDERR
E0000 00:00:1742438694.353348      391 cuda_dnn.cc:8579] Unable to
register cuDNN factory: Attempting to register factory for plugin
cuDNN when one has already been registered
E0000 00:00:1742438694.362098      391 cuda_blas.cc:1407] Unable to
register cuBLAS factory: Attempting to register factory for plugin
cuBLAS when one has already been registered
W0000 00:00:1742438694.406488      391 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1742438694.406553      391 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1742438694.406555      391 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
W0000 00:00:1742438694.406557      391 computation_placer.cc:177]
computation placer already registered. Please check linkage and avoid
linking the same target more than once.
2025-03-19 19:44:54.414449: I
tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow

```

binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

2.19.0

```
file_id = "1oQ6Vy_HqZlVHnkFspgxMn0IcE__D8Kmh"  
zip_filename = "ocular-disease-recognition.zip"  
extract_path = "./ocular-disease-recognition"
```

```
# Check if the file already exists
```

```
if not os.path.exists(zip_filename):  
    print(f"Downloading {zip_filename}...")  
    gdown.download(f"https://drive.google.com/uc?id={file_id}",  
zip_filename, quiet=False)  
else:  
    print(f"{zip_filename} already exists. Skipping download.")
```

ocular-disease-recognition.zip already exists. Skipping download.

```
# Check if already extracted
```

```
if not os.path.exists(extract_path):  
    os.makedirs(extract_path, exist_ok=True)  
    print(f"Extracting {zip_filename}...")  
  
    with zipfile.ZipFile(zip_filename, "r") as zip_ref:  
        zip_ref.extractall(extract_path)  
  
    print(f"Extraction complete! Files extracted to: {extract_path}")  
else:  
    print(f"Extraction skipped: {extract_path} already exists.")
```

Extraction skipped: ./ocular-disease-recognition already exists.

```
# Load the dataset (Update the path if necessary)
```

```
dataset_path = "processed_ocular_disease.csv"  
df = pd.read_csv(dataset_path)
```

```
df = df[  
    df.apply(lambda row:
```

```
os.path.exists(os.path.join('ocular-disease-recognition/preprocessed_i  
mages', row['Left-Fundus'])) and
```

```
os.path.exists(os.path.join('ocular-disease-recognition/preprocessed_i  
mages', row['Right-Fundus'])),  
    axis=1  
    )
```

```
].reset_index(drop=True)
```

```

# **Split into Train, Validation, and Test Sets**
train_df, temp_df = train_test_split(df, test_size=0.3,
random_state=42) # 70% Train, 30% Temp
val_df, test_df = train_test_split(temp_df, test_size=0.5,
random_state=42) # 15% Val, 15% Test

# **Get Number of Classes**
num_classes = len(np.unique(df['labels']))
print(f"Number of Classes: {num_classes}")

# Compute class weights
class_labels = np.unique(df['labels'])
class_weights = compute_class_weight(class_weight="balanced",
classes=class_labels, y=df['labels'])
class_weight_dict = {i: class_weights[i] for i in
range(len(class_labels))}

print("Computed Class Weights:", class_weight_dict)

Number of Classes: 8
Computed Class Weights: {0: np.float64(2.917307692307692), 1:
np.float64(2.718637992831541), 2: np.float64(0.49285250162443145), 3:
np.float64(2.7783882783882783), 4: np.float64(6.01984126984127), 5:
np.float64(3.77363184079602), 6: np.float64(0.2791682002208318), 7:
np.float64(1.1270430906389302)}

from sklearn.utils import resample

def balance_classes(df):
    """Resample dataset to balance classes."""
    max_size = df['labels'].value_counts().max() # Find the max
number of samples in any class
    balanced_df = pd.concat([
        resample(df[df['labels'] == cls], replace=True,
n_samples=max_size, random_state=42)
        for cls in df['labels'].unique()
    ])
    return balanced_df.sample(frac=1).reset_index(drop=True) #
Shuffle after resampling

# Apply to training data only
train_df_balanced = balance_classes(train_df)

# Check if balancing worked
print(train_df_balanced['labels'].value_counts()) # Should now be
balanced

labels
3    1901
0    1901
6    1901

```

```
2    1901
7    1901
4    1901
5    1901
1    1901
```

```
Name: count, dtype: int64
```

```
class DualImageAugmentation(DualTransform):
    def __init__(self, transforms, always_apply=False, p=0.5):
        super(DualImageAugmentation, self).__init__(always_apply, p)
        self.transforms = A.Compose(transforms)

    def apply(self, img, **params):
        return self.transforms(image=img)["image"]

    def apply_to_image1(self, img, **params):
        return self.transforms(image=img)["image"]

class OcularDatasetGenerator(Sequence):
    def __init__(self, df, batch_size=32, img_size=(128, 128),
shuffle=True, augment=True, **kwargs):
        super().__init__(**kwargs)

        self.df = df[df.apply(lambda row:
os.path.exists(os.path.join('ocular-disease-recognition/preprocessed_i
mages', row['Left-Fundus'])) and

os.path.exists(os.path.join('ocular-disease-recognition/preprocessed_i
mages', row['Right-Fundus']))),
axis=1
)].reset_index(drop=True) # Reset index after filtering

        print(f"Dataset initialized with {len(self.df)} valid
samples.")

        self.batch_size = batch_size
        self.img_size = img_size
        self.shuffle = shuffle
        self.augment = augment
        self.indices = np.arange(len(df))

        # Define augmentation pipeline if augmentation is enabled
        if augment:
            self.augmentation_pipeline =
self.get_augmentation_pipeline()
        else:
            self.augmentation_pipeline = None
```

```

        self.on_epoch_end()

    def __len__(self):
        return int(np.floor(len(self.df) / self.batch_size))

    def __getitem__(self, index):
        batch_indices = self.indices[index * self.batch_size:(index +
1) * self.batch_size]
        batch = self.df.iloc[batch_indices]
        X, y = self.__data_generation(batch)
        return np.array(X), np.array(y)

    # def __data_generation(self, batch):
    #     X_batch = []
    #     y_batch = []

    #     for _, row in batch.iterrows():
    #         left_image_path = os.path.join('ocular-disease-
recognition/preprocessed_images', row['Left-Fundus'])
    #         right_image_path = os.path.join('ocular-disease-
recognition/preprocessed_images', row['Right-Fundus'])

    #         left_image = self.load_image(left_image_path)
    #         right_image = self.load_image(right_image_path)

    #         if left_image is None or right_image is None:
    #             continue

    #         # Apply identical augmentations to both images
    #         if self.augment and self.augmentation_pipeline:
    #             augmented =
self.augmentation_pipeline(image=left_image, image1=right_image)
    #             left_image = augmented["image"]
    #             right_image = augmented["image1"]

    #         # Ensure channel dimension is included (for grayscale
images)
    #         left_image = np.expand_dims(left_image, axis=-1) #
Shape: (128, 128, 1)
    #         right_image = np.expand_dims(right_image, axis=-1) #
Shape: (128, 128, 1)

    #         combined_image = np.concatenate((left_image,
right_image), axis=-1) # Shape: (128, 128, 2)
    #         X_batch.append(combined_image)

    #         y_batch.append(int(row['labels']))

    #     return np.array(X_batch, dtype=np.float32),
np.array(y_batch, dtype=np.int32)

```

```

def __data_generation(self, batch):
    X_batch = []
    y_batch = []

    for _, row in batch.iterrows():
        left_image_path = os.path.join('ocular-disease-
recognition/preprocessed_images', row['Left-Fundus'])
        right_image_path = os.path.join('ocular-disease-
recognition/preprocessed_images', row['Right-Fundus'])

        left_image = self.load_image(left_image_path)
        right_image = self.load_image(right_image_path)

        if left_image is None or right_image is None:
            continue # Skip invalid images

        # Apply augmentation (both images get the same
transformation)
        if self.augment and self.augmentation_pipeline:
            augmented =
self.augmentation_pipeline(image=left_image, image1=right_image)
            left_image = augmented["image"]
            right_image = augmented["image1"]

        # Convert grayscale images to 3D (required for CNN)
        left_image = np.expand_dims(left_image, axis=-1)
        right_image = np.expand_dims(right_image, axis=-1)

        # Merge images into a two-channel input
        combined_image = np.concatenate((left_image, right_image),
axis=-1)

        X_batch.append(combined_image)
        y_batch.append(int(row['labels']))

    return np.array(X_batch, dtype=np.float32), np.array(y_batch,
dtype=np.int32)

def load_image(self, image_path):
    image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
    if image is None:
        return None
    image = cv2.resize(image, self.img_size)
    image = image / 255.0
    return image

# def load_image(self, image_path):

```

```

#     image = cv2.imread(image_path, cv2.IMREAD_GRAYSCALE)
#     if image is None:
#         return None

#     image = cv2.resize(image, self.img_size)

#     # □ Ensure image is uint8 before passing to CLAHE
#     image = np.clip(image, 0, 255).astype(np.uint8)

#     return image

# def get_augmentation_pipeline(self):
#     return DualImageAugmentation([
#         A.RandomBrightnessContrast(p=0.4),
#         A.GaussianBlur(blur_limit=(3, 7), p=0.3),
#         A.HorizontalFlip(p=0.5),
#         A.Affine(rotate=(-20, 20), scale=(0.9, 1.1),
translate_percent=(0.05, 0.05), p=0.6),
#         A.ElasticTransform(p=0.3),
#         A.GridDistortion(p=0.3),
#         A.RandomRotate90(p=0.3),
#         A.CLAHE(p=0.2), # Contrast enhancement
#         A.Cutout(num_holes=1, max_h_size=20, max_w_size=20,
p=0.5), # □ Helps prevent over-reliance on features
#     ])

# def get_augmentation_pipeline(self):
#     return A.Compose([
#         A.RandomBrightnessContrast(p=0.5),
#         A.GaussianBlur(blur_limit=(3, 7), p=0.4),
#         A.HorizontalFlip(p=0.5),
#         A.Rotate(limit=30, p=0.5), # Add rotations
#         A.ElasticTransform(p=0.5), # More distortions
#         A.CoarseDropout(num_holes_range=(1, 3),
hole_height_range=(0.05, 0.2), hole_width_range=(0.05, 0.2), p=0.5),
#     ])

def get_augmentation_pipeline(self):
    return A.Compose([
        A.RandomBrightnessContrast(p=0.5),
        A.GaussianBlur(blur_limit=(3, 7), p=0.4),
        A.HorizontalFlip(p=0.5),
        A.Rotate(limit=30, p=0.5),
        A.ElasticTransform(p=0.5),
        A.CoarseDropout(max_holes=3, max_height=0.2,
max_width=0.2, p=0.5),
    ], additional_targets={"image1": "image"})

```

```

def on_epoch_end(self):
    """ Shuffle indices at the end of each epoch. """
    self.indices = np.arange(len(self.df)) # Ensure indices match
    filtered dataset
    if self.shuffle:
        np.random.shuffle(self.indices)

# **Create Data Generators**
batch_size = 64
train_generator = OcularDatasetGenerator(train_df_balanced,
batch_size=batch_size, img_size=(224, 224), augment=True)
val_generator = OcularDatasetGenerator(val_df, batch_size=batch_size,
img_size=(224, 224))
test_generator = OcularDatasetGenerator(test_df,
batch_size=batch_size, shuffle=False, img_size=(224, 224))

# **Define CNN Model**
model = models.Sequential([
    layers.Input(shape=(224, 224, 2)),

    layers.Conv2D(32, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.1), # Reduced early dropout

    layers.Conv2D(64, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.2), # Keep moderate dropout

    layers.Conv2D(128, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.2),

    layers.Conv2D(256, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.3),

    layers.Conv2D(512, (3, 3), activation='relu', padding='same'),
    layers.BatchNormalization(),
    layers.MaxPooling2D((2, 2)),
    layers.Dropout(0.4),

    layers.Flatten(),
    layers.Dense(512, activation='relu'),
    layers.BatchNormalization(),

```



```

        layers.Dropout(0.5), # Keep high dropout only in the final
dense layer

        layers.Dense(num_classes, activation='softmax')
    ])

def step_decay(epoch):
    initial_lr = 0.001
    drop = 0.5
    epochs_drop = 5
    return initial_lr * (drop ** (epoch // epochs_drop))

def cosine_decay(epoch):
    initial_lr = 0.001
    return initial_lr * (0.5 * (1 + np.cos(np.pi * epoch / 100))) #
100 = total epochs

lr_schedule = tf.keras.callbacks.LearningRateScheduler(cosine_decay)

model.compile(
    optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),
    loss='sparse_categorical_crossentropy', # Use sparse version
    metrics=['accuracy']
)

# **Define Early Stopping Callback**
early_stopping = callbacks.EarlyStopping(
    monitor='val_loss', # Stop if validation loss stops improving
    patience=8, # Wait for 5 epochs before stopping
    restore_best_weights=True # Restore best model weights
)

# **Train the Model with Early Stopping**
history = model.fit(
    train_generator,
    epochs=100,
    validation_data=val_generator,
    callbacks=[early_stopping, lr_schedule],
)

```

Dataset initialized with 15208 valid samples.

Dataset initialized with 910 valid samples.

Dataset initialized with 911 valid samples.

/tmp/ipykernel_391/3805339739.py:165: UserWarning: Argument(s)
'max_holes, max_height, max_width' are not valid for transform
CoarseDropout

A.CoarseDropout(max_holes=3, max_height=0.2, max_width=0.2, p=0.5),
I0000 00:00:1742438785.377546 391 gpu_device.cc:2019] Created

```
device /job:localhost/replica:0/task:0/device:GPU:0 with 8847 MB
memory: -> device: 0, name: NVIDIA GeForce RTX 3080, pci bus id:
0000:06:00.0, compute capability: 8.6
```

Epoch 1/100

```
WARNING: All log messages before absl::InitializeLog() is called are
written to STDERR
I0000 00:00:1742438793.306045      752 service.cc:152] XLA service
0x7f74a4004d20 initialized for platform CUDA (this does not guarantee
that XLA will be used). Devices:
I0000 00:00:1742438793.306183      752 service.cc:160] StreamExecutor
device (0): NVIDIA GeForce RTX 3080, Compute Capability 8.6
2025-03-19 19:46:33.430513: I
tensorflow/compiler/mlir/tensorflow/utils/dump_mlir_util.cc:269]
disabling MLIR crash reproducer, set env var
`MLIR_CRASH_REPRODUCER_DIRECTORY` to enable.
I0000 00:00:1742438794.097146      752 cuda_dnn.cc:529] Loaded cuDNN
version 90300
2025-03-19 19:46:35.293984: I
external/local_xla/xla/stream_executor/cuda/subprocess_compilation.cc:
346] ptxas warning : Registers are spilled to local memory in function
'gemm_fusion_dot_4658', 4 bytes spill stores, 4 bytes spill loads
```

```
1/237 _____ 1:07:55 17s/step - accuracy: 0.0938 -
loss: 3.3140
```

```
I0000 00:00:1742438806.206409      752 device_compiler.h:188] Compiled
cluster using XLA! This line is logged at most once for the lifetime
of the process.
```

```
237/237 _____ 77s 254ms/step - accuracy: 0.2376 - loss:
2.4636 - val_accuracy: 0.0949 - val_loss: 2.7261 - learning_rate:
0.0010
```

Epoch 2/100

```
237/237 _____ 56s 235ms/step - accuracy: 0.3270 - loss:
1.9466 - val_accuracy: 0.0525 - val_loss: 3.6943 - learning_rate:
9.9975e-04
```

Epoch 3/100

```
237/237 _____ 54s 226ms/step - accuracy: 0.3961 - loss:
1.6655 - val_accuracy: 0.1908 - val_loss: 1.9448 - learning_rate:
9.9901e-04
```

Epoch 4/100

```
237/237 _____ 56s 237ms/step - accuracy: 0.4660 - loss:
1.4428 - val_accuracy: 0.2377 - val_loss: 1.7995 - learning_rate:
9.9778e-04
```

Epoch 5/100

```
237/237 _____ 52s 220ms/step - accuracy: 0.5132 - loss:
1.3165 - val_accuracy: 0.2589 - val_loss: 1.7174 - learning_rate:
```

9.9606e-04
Epoch 6/100
237/237 _____ 52s 217ms/step - accuracy: 0.5481 - loss: 1.2080 - val_accuracy: 0.3147 - val_loss: 1.6692 - learning_rate: 9.9384e-04
Epoch 7/100
237/237 _____ 52s 217ms/step - accuracy: 0.5937 - loss: 1.0974 - val_accuracy: 0.2924 - val_loss: 1.7421 - learning_rate: 9.9114e-04
Epoch 8/100
237/237 _____ 52s 220ms/step - accuracy: 0.6170 - loss: 1.0391 - val_accuracy: 0.2455 - val_loss: 1.9066 - learning_rate: 9.8796e-04
Epoch 9/100
237/237 _____ 52s 219ms/step - accuracy: 0.6100 - loss: 1.0706 - val_accuracy: 0.3013 - val_loss: 1.6941 - learning_rate: 9.8429e-04
Epoch 10/100
237/237 _____ 55s 232ms/step - accuracy: 0.6457 - loss: 0.9456 - val_accuracy: 0.3359 - val_loss: 1.5632 - learning_rate: 9.8015e-04
Epoch 11/100
237/237 _____ 63s 265ms/step - accuracy: 0.6786 - loss: 0.8730 - val_accuracy: 0.3170 - val_loss: 1.6493 - learning_rate: 9.7553e-04
Epoch 12/100
237/237 _____ 60s 252ms/step - accuracy: 0.7170 - loss: 0.7669 - val_accuracy: 0.4007 - val_loss: 1.4947 - learning_rate: 9.7044e-04
Epoch 13/100
237/237 _____ 53s 223ms/step - accuracy: 0.7205 - loss: 0.7412 - val_accuracy: 0.3917 - val_loss: 1.5634 - learning_rate: 9.6489e-04
Epoch 14/100
237/237 _____ 52s 221ms/step - accuracy: 0.7232 - loss: 0.7505 - val_accuracy: 0.4208 - val_loss: 1.4743 - learning_rate: 9.5888e-04
Epoch 15/100
237/237 _____ 52s 219ms/step - accuracy: 0.7336 - loss: 0.7126 - val_accuracy: 0.3962 - val_loss: 1.6253 - learning_rate: 9.5241e-04
Epoch 16/100
237/237 _____ 52s 220ms/step - accuracy: 0.7625 - loss: 0.6419 - val_accuracy: 0.3951 - val_loss: 1.5768 - learning_rate: 9.4550e-04
Epoch 17/100
237/237 _____ 52s 220ms/step - accuracy: 0.7524 - loss: 0.6642 - val_accuracy: 0.4475 - val_loss: 1.5106 - learning_rate: 9.3815e-04

```
Epoch 18/100
237/237 _____ 52s 219ms/step - accuracy: 0.7656 - loss:
0.6384 - val_accuracy: 0.4118 - val_loss: 1.5936 - learning_rate:
9.3037e-04
Epoch 19/100
237/237 _____ 53s 223ms/step - accuracy: 0.7845 - loss:
0.5806 - val_accuracy: 0.3973 - val_loss: 1.6955 - learning_rate:
9.2216e-04
Epoch 20/100
237/237 _____ 51s 217ms/step - accuracy: 0.7917 - loss:
0.5803 - val_accuracy: 0.4420 - val_loss: 1.5719 - learning_rate:
9.1354e-04
Epoch 21/100
237/237 _____ 52s 218ms/step - accuracy: 0.8120 - loss:
0.5151 - val_accuracy: 0.4531 - val_loss: 1.5709 - learning_rate:
9.0451e-04
Epoch 22/100
237/237 _____ 51s 216ms/step - accuracy: 0.8060 - loss:
0.5348 - val_accuracy: 0.4754 - val_loss: 1.4930 - learning_rate:
8.9508e-04

# **Evaluate the Model on Test Set**
test_loss, test_acc = model.evaluate(test_generator)
print(f"Test Accuracy: {test_acc:.4f}")

14/14 _____ 3s 217ms/step - accuracy: 0.4259 - loss:
1.4151
Test Accuracy: 0.4241
```